




UNIVERSITEIT VAN AMSTERDAM

# N-Body Simulation

Barnes-Hut Approximation

# Introduction: N-Body Simulation



- Used to predict future particle (body) positions
- Each body has:
  - Mass
  - Position
  - Velocity
- The simulation runs for a given number of iterations
- Each iteration:
  - Body positions are updated based on velocities
  - Body velocities are updated by calculating the forces acting on the bodies using Newton's formulas
- By nature  $O(n^2)$
- Goal: Optimize & Model for execution time



# Introduction: Barnes-Hut Approximation

- Estimated to be  $O(n \log n)$
- Produces a quadtree based on the positions of bodies
- Approximates distant body forces by grouping distant bodies
  - Approximation occurs when:  $\frac{diag}{dist} < \theta$
- Aiming for a maximum error of less than 1% after 50 iterations

Theta	Error
0.01	0%
0.05	0.06%
0.1	0.13%
0.5	1.05%
1	3.94%
1.5	11.74%



# Experimental Setup

## Software

### Universe configuration

<b>Bodies</b> (32-bit unsigned int)	[1e2, 1e4]
<b>Mass</b> [kg] (64-bit float)	[1e10, 1e40]
<b>Velocity-X</b> [m] (64-bit float)	[-1e6, 1e6]
<b>Velocity-Y</b> [m] (64-bit float)	[-1e6, 1e6]
<b>Position-X</b> [m/s] (64-bit float)	[-5e16, 5e16]
<b>Position-Y</b> [m/s] (64-bit float)	[-5e16, 5e16]

### Simulation configuration

<b>Iterations</b> (32-bit unsigned int)	50
<b>Iteration Duration</b> [s] (64-bit float)	3600
<b>Theta</b> (Barnes-Hut only) (64-bit float)	0.5

## Hardware

### Intel® Xeon® E5-2630 v3 Processor

Physical cores	8
Logical cores	16
Base frequency	2.40 GHz
Turbo frequency	3.20 GHz
Size L1 cache	32 KB
Size L2 cache	256 KB
Size L3 cache	20 MB
Latency L1 cache	4 cycles
Latency L2 cache	12 cycles
Latency L3 cache	34 cycles
Cache line size	64 Bytes



# Prototype 1: Model

Symbol	Description
$t_{pu}$	Time required to update the position of a single body
$t_{tl}$	Time required to insert a body in a quads subtree
$t_q$	Time required to generate empty subtree
$t_{rc}$	Time required to call a function
$t_{fc}$	Time required to calculate a force acting on a body
$t_{dc}$	Time required to calculate the Euclidean distance between any two points
$Dp(n)$	Estimator of average tree depth
$Fc(n)$	Estimator of number of force calculations per body
$Nv(n)$	Estimator of number of node visits per body
$Lvf(n)$	Estimator of ratio: leaf_visits / total_visits

$$T = T_{pu} + T_{tg} + T_{vu}$$

$$T_{pu} = N \cdot t_{pu}$$


$$T_{tg} = N \cdot \left( Dp(N) \cdot t_{tl} + \frac{4(t_q + t_{rc})}{3} \right)$$

$$T_{vu} = N \cdot \left( Fc(N) \cdot t_{fc} + Nv(N) \cdot t_{rc} + (1 - Lvf(N)) \cdot Nv(N) \cdot t_{dc} \right)$$



# Prototype 1: Approximation Functions

- Understand the nature of the phenomena
- Gather measurement on a training set
- Train the model
- Validate it on a test set
- Enjoy your approximation



# Prototype 1: Model Calibration

Symbol	Value
$t_{pu}$	5ns
$t_{tl}$	29ns
$t_q$	60ns
$t_{rc}$	0.6ns
$t_{fc}$	55ns
$t_{dc}$	34ns
$Dp(n)$	$1.64 \cdot \ln(n) - 0.91$
$Fc(n)$	$45 \cdot \ln(n) - 160$
$Nv(n)$	$65 \cdot \ln(n) - 197$
$Lvf(n)$	$-0.0879 \cdot \ln(n) + 1.13$

$$T = N \cdot (194.25 \cdot \ln^2(N) + 1,685.50 \cdot \ln(N) - 7,988.05)$$



# Prototype 1: Results

<b>N</b>	<b>Measured[s]</b>	<b>Predicted[s]</b>	<b>Error</b>
10000	0.2995760	0.2403889	19.76%
5000	0.1132730	0.1024209	9.58%
1000	0.01404501	0.0129514	7.79%
500	0.00545305	0.0050089	8.15%
100	0.00033501	0.0003927	17.22%

<b>N</b>	<b>T<sub>pu</sub></b>	<b>T<sub>tg</sub></b>	<b>T<sub>vu</sub></b>
10000	0.01%	1.77%	98.23%
5000	0.01%	2.17%	97.82%
1000	0.02%	3.41%	96.57%
500	0.03%	4.42%	95.55%
100	0.15%	15.10%	84.76%





# Prototype 2: Optimisations

- Square root and division computations are expensive:

- $\frac{diag}{dist} < \theta \longrightarrow diag^2 < \theta^2 \cdot dist^2$

- `#pragma omp parallel for schedule(dynamic)`

- Applied on the *Velocity Update* loop
  - Dynamic scheduling determined optimal due to load imbalance



## Prototype 2: Model

$$T = T_{pu} + T_{tg} + T_{vu}$$

$$T_{pu} = N \cdot t_{pu}$$

$$T_{tg} = N \cdot \left( Dp(N) \cdot t_{tl} + \frac{4(t_q + t_{rc})}{3} \right)$$

$$T_{vu} = \frac{N}{P} \cdot \left( Fc(N) \cdot Tfc(P) + Nv(N) \cdot t_{rc} + (1 - Lv f(N)) \cdot Nv(N) \cdot Tdc(P) \right)$$

Symbol	Description
$Tfc(p)$	Estimator of time required to calculate a force acting on a body
$Tdc(p)$	Estimator of time required to calculate the Euclidean distance between any two points



## Prototype 2: Model Calibration

Symbol	Value
$t_{pu}$	8ns
$t_{tl}$	68ns
$t_q$	105ns
$t_{rc}$	0.6ns
$Tfc(p)$	$2.23 \cdot p + 52.1$
$Tdc(p)$	$2.27 \cdot p + 28.5$
$Dp(n)$	$1.64 \cdot \ln(n) - 0.91$
$Fc(n)$	$45 \cdot \ln(n) - 160$
$Nv(n)$	$65 \cdot \ln(n) - 197$
$Lvf(n)$	$-0.0879 \cdot \ln(n) + 1.13$

$$T = N \cdot (12.96 \cdot \ln^2(N) + 153.38 \cdot \ln(N) - 211.74 + \frac{162.83 \cdot \ln^2(N) + 1649.16 \cdot \ln(N) - 7724.31}{P})$$



## Prototype 2: Results

N	P	Measured[s]	Predicted[s]	Error
10000	2	0.104973	0.129726	23.58%
10000	4	0.066044	0.076354	15.61%
10000	8	0.040812	0.049668	21.70%
1000	2	0.007450	0.007214	3.16%
1000	4	0.004567	0.004339	4.98%
1000	8	0.003095	0.002902	6.23%
100	2	0.000495	0.000246	50.35%
100	4	0.000364	0.000161	55.64%
100	8	0.000343	0.000119	65.29%

N	P	T <sub>pu</sub>	T <sub>tg</sub>	T <sub>vu</sub>
10000	2	0.07%	10.49%	89.44%
10000	4	0.12%	17.04%	82.85%
10000	8	0.20%	28.26%	71.54%
1000	2	0.14%	14.94%	84.92%
1000	4	0.23%	24.21%	75.55%
1000	8	0.38%	38.10%	61.52%
100	2	0.28%	21.32%	78.40%
100	4	0.37%	29.49%	70.14%
100	8	0.31%	32.78%	66.91%



# Prototype 3: Optimisations

- Bulk quads initialisation
  - Less system calls
  - Better caching as quadtree is stored in an array
- Replacing **std::list** with **std::vector**
  - Better caching behaviour, faster iterations
  - Allocated with initial size of 40



## Prototype 3: Model

$$T = T_{pu} + T_{tg} + T_{vu}$$

$$T_{pu} = N \cdot t_{pu}$$

$$T_{tg} = N \cdot \left( Dp(N) \cdot t_{tl} + 4(t_q + t_{rc}) \right)$$

$$T_{vu} = \frac{N}{P} \cdot \left( Fc(N) \cdot Tfc(P) + Nv(N) \cdot t_{rc} + (1 - Lvf(N)) \cdot Nv(N) \cdot Tdc(P) \right)$$



# Prototype 3: Model Calibration

Symbol	Value
$t_{pu}$	8ns
$t_{tl}$	11ns
$t_q$	80ns
$t_{rc}$	0.6ns
$Tfc(p)$	$2.16 \cdot p + 35.6$
$Tdc(p)$	$2.27 \cdot p + 28.5$
$Dp(n)$	$1.64 \cdot \ln(n) - 0.91$
$Fc(n)$	$45 \cdot \ln(n) - 160$
$Nv(n)$	$65 \cdot \ln(n) - 197$
$Lvf(n)$	$-0.0879 \cdot \ln(n) + 1.13$

$$T = N \cdot (12.96 \cdot \ln^2(N) + 56.75 \cdot \ln(N) + 32.92 + \frac{162.83 \cdot \ln^2(N) + 947.16 \cdot \ln(N) - 5228.31}{P})$$



# Prototype 3: Results

N	P	Measured[s]	Predicted[s]	Error
10000	2	0.090472	0.102891	13.73%
10000	4	0.049679	0.059738	20.25%
10000	8	0.030900	0.038161	23.50%
1000	2	0.006213	0.005581	10.18%
1000	4	0.003907	0.003313	15.20%
1000	8	0.002480	0.002180	12.11%
100	2	0.000325	0.000188	42.31%
100	4	0.000284	0.000122	56.91%
100	8	0.000250	0.000090	64.12%

N	P	Tpu	Ttg	Tvu
10000	2	0.09%	5.19%	94.72%
10000	4	0.17%	8.88%	90.95%
10000	8	0.28%	14.30%	85.43%
1000	2	0.16%	9.20%	90.64%
1000	4	0.28%	14.37%	85.35%
1000	8	0.47%	22.76%	76.77%
100	2	0.44%	21.91%	77.65%
100	4	0.49%	26.52%	72.99%
100	8	0.59%	28.64%	70.76%





# Calibrated Models: Overview

## Model 1:

$$T = N \cdot (194.25 \cdot \ln^2(N) + 1,685.50 \cdot \ln(N) - 7,988.05)$$

## Model 2:

$$T = N \cdot (12.96 \cdot \ln^2(N) + 153.38 \cdot \ln(N) - 211.74 + \frac{162.83 \cdot \ln^2(N) + 1649.16 \cdot \ln(N) - 7724.31}{P})$$


## Model 3:

$$T = N \cdot (12.96 \cdot \ln^2(N) + 56.75 \cdot \ln(N) + 32.92 + \frac{162.83 \cdot \ln^2(N) + 947.16 \cdot \ln(N) - 5228.31}{P})$$




# Performance Prediction

**2 threads**

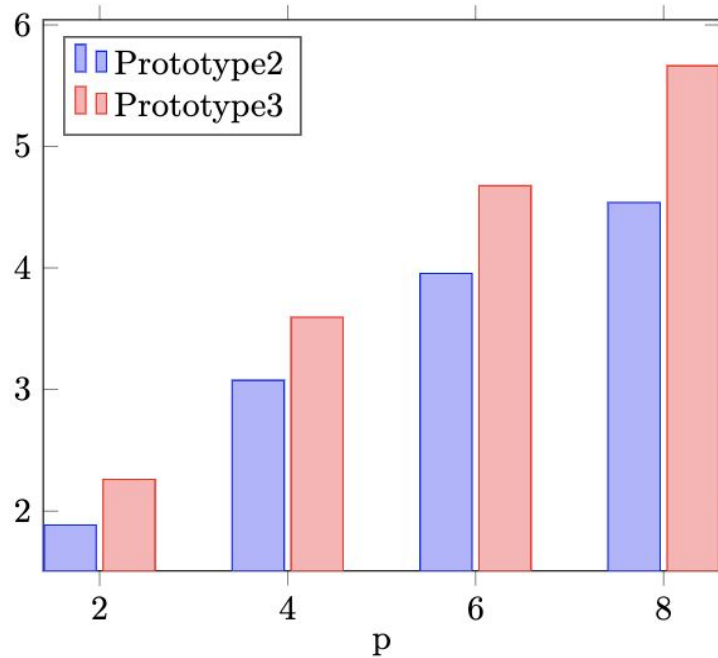


**8 threads**

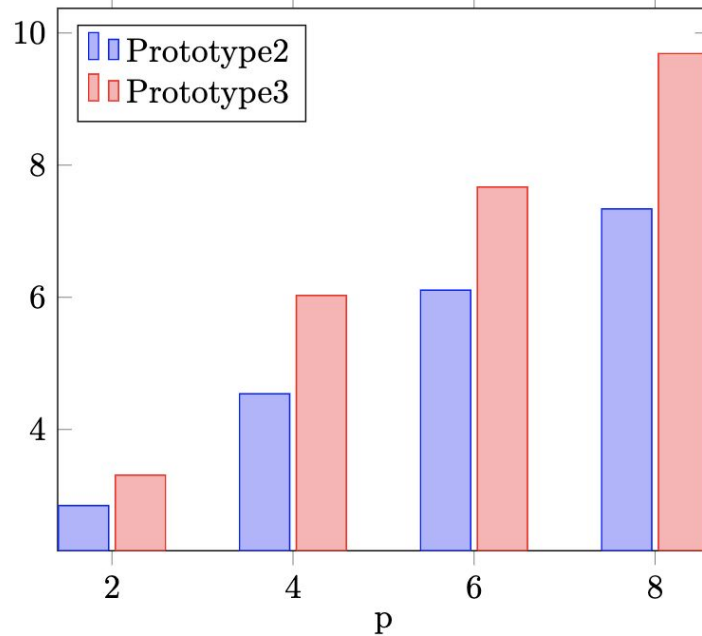


# Performance Results

1k particles



10k particles



# Final Result

**1.89s -> 0.0376s**

**Speedup: 50.4**

